

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from datetime import datetime
from datetime import timedelta
from pandas.plotting import register_matplotlib_converters
from statsmodels.tsa.stattools import acf, pacf
from statsmodels.tsa.arima.model import ARIMA
register_matplotlib_converters()
from time import time
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf, month_plot, quarter_plot

from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
In [13]: df_ice_cream = pd.read_csv("ice_cream.csv")
df = df.copy()
df_ice_cream.head()
```

Out[13]:

	DATE	IPN31152N
0	1972-01-01	59.9622
1	1972-02-01	67.0605
2	1972-03-01	74.2350
3	1972-04-01	78.1120
4	1972-05-01	84.7636

```
In [15]: df_ice_cream = df_ice_cream.rename(columns = {'DATE':'date','IPN31152N':'production'})

In [17]: df_ice_cream['date'] = pd.to_datetime(df_ice_cream.date)

In [19]: df_ice_cream.set_index('date',inplace=True)


In [21]: start_date = pd.to_datetime('2010-01-01')
df_ice_cream = df_ice_cream[df_ice_cream.index >= start_date]

In [23]: df_ice_cream.head()
```


Out[23]:

	production
date	
2010-01-01	91.2895
2010-02-01	110.4994
2010-03-01	127.0971
2010-04-01	132.6468
2010-05-01	134.5576

```
In [25]: plt.figure(figsize=(10,4))
plt.plot(df_ice_cream.production)
plt.title('Ice Cream Production over Time', fontsize=20)
plt.ylabel('Production', fontsize=16)
for year in range(2010,2021):
    plt.axvline(pd.to_datetime(str(year)+'-01-01'), color='k', linestyle='--', alpha=0.2)
```



```
In [27]: acf_plot = plot_acf(df_ice_cream.production,lags=100)
```



```
In [29]: train_end = datetime(2018,12,1)
test_end = datetime(2019,12,1)

train_data = df_ice_cream.production[:train_end]
test_data = df_ice_cream.production[train_end + timedelta(days=1):test_end]
```

```
In [31]: from statsmodels.tsa.arima.model import ARIMA
model = ARIMA(train_data, order=(3,0,0))

C:\Users\LEGION\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
self._init_dates(dates, freq)
C:\Users\LEGION\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
self._init_dates(dates, freq)
C:\Users\LEGION\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
self._init_dates(dates, freq)
```

```
In [33]: start = time()
model_fit = model.fit()
end = time()
print('Model fitting time:',end-start)

Model fitting time: 0.16944384574890137
```

```
In [35]: print(model_fit.summary())


SARIMAX Results
=====
Dep. Variable:      production    No. Observations:      108
Model:              ARIMA(3, 0, 0)    Log Likelihood      -374.085
Date:               Wed, 04 Sep 2024    AIC                  758.170
Time:               20:40:56           BIC                  771.580
Sample:             01-01-2010        HQIC                 763.607
                             - 12-01-2018
Covariance Type:    opg
=====
                coef      std err          2    P>|z|          2    [0.025    0.975]
-----
const      103.5743      2.169      47.761    0.000      99.324    107.825
ar.L1       1.0469      0.102     10.256    0.000      0.847      1.247
ar.L2      -0.0523      0.175      -0.298    0.765     -0.396      0.291
ar.L3      -0.4044      0.123      -3.290    0.001     -0.645     -0.164
sigma2      58.4026      9.475      6.164    0.000      39.831      76.974
=====
Ljung-Box (L1) (Q):      0.63    Jarque-Bera (JB):      8.71
Prob(Q):                0.43    Prob(JB):              0.01
Heteroskedasticity (H):  0.76    Skew:                0.70
Prob(H) (two-sided):    0.41    Kurtosis:           3.02
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

```
In [37]: pred_start_date = test_data.index[0]
pred_end_date = test_data.index[-1]
```


```
In [41]: predictions = model_fit.predict(start=pred_start_date, end = pred_end_date)
residuals = test_data - predictions

plt.figure(figsize=(10,4))
plt.plot(residuals)
plt.title('Residuals from AR model',fontsize=20)
plt.ylabel('Error',fontsize=16)
plt.axhline(0,color='r',linestyle='--',alpha=0.2)
for year in range(2019,2021):
    plt.axvline(pd.to_datetime(str(year)+'-01-01'),color='k',linestyle='--',alpha=0.2)
```



```
In [43]: plt.figure(figsize=(10,4))
plt.plot(test_data)
plt.plot(predictions)
plt.legend(['Data','Predictions'],fontsize=16)
plt.title('Ice cream productions over time',fontsize=20)
plt.ylabel('Production',fontsize=16)
```

```
Out[43]: Text(0, 0.5, 'Production')
```



```
In [45]: print('Mean Absolute Percent Error:', round(np.mean(abs(residuals/test_data)),4))
print('Root mean squared error:', np.sqrt(np.mean(residuals**2)))
```

Mean Absolute Percent Error: 0.0895  
Root mean squared error: 9.88415874704018

In [ ]: